

Graphs in Manufacturing

mark.quinsland@neo4j.com

May 24, 2018



Neo4j - The Graph Company

Industry's Largest Dedicated Investment in Graphs



- Creator of the Neo4j Graph Platform
- ~200 employees
- HQ in Silicon Valley, other offices include London, Munich, Paris and Malmö (Sweden)
- \$80M in funding from Fidelity, Sunstone, Conor, Creandum, and Greenbridge Capital
- Over 10M+ downloads
- 270+ enterprise subscription customers with over half with >\$1B in revenue

Adoption

- 7/10 Top Retail Firms
- 12/25 Top Financial Firms
- 8/10 Top Software Vendors

Ecosystem

- 720+ Startups in program
- 270+ Enterprise customers
- 100+ Partners
- 53K+ Meet up members
- 450+ Events per year

Customers



COMCAST



Lufthansa



Partners



pitney bowes



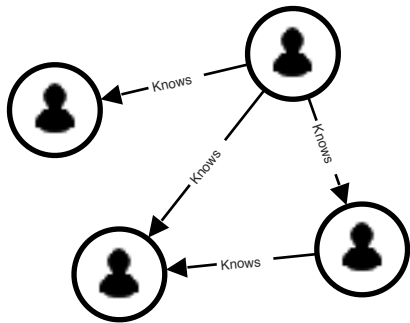
GraphAware

structr

Why Graph?


The Rise of Connections in Data

Data connections are increasing as rapidly as data volumes




Networks of People

E.g., Employees,
Customers, Suppliers,
Partners, Influencers



Business Processes

E.g., Risk management,
Supply chain, Payments

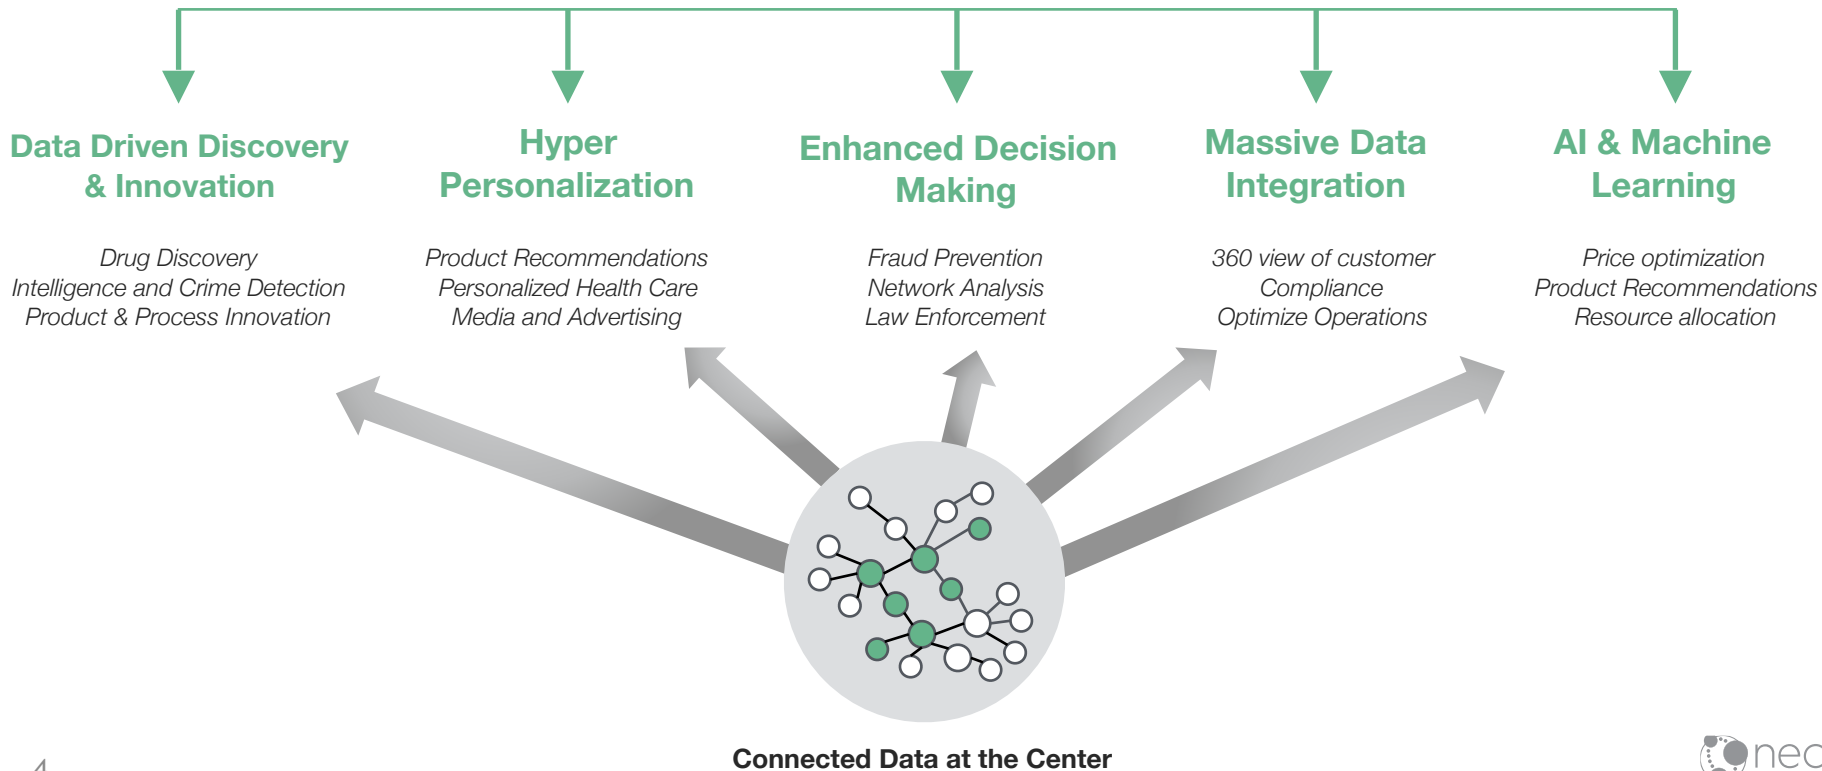


Knowledge Networks

E.g., Enterprise content,
Domain specific content,
eCommerce content

Harnessing Connections Drives Business Value

Digital Transformation Megatrends



Graphs Are a Logical Choice for Many Areas



Organization




Figure 5-7. Talent.net graph enriched with MARKED_WITH relationships

Identity & Access




Figure 5-8. Access control graph

Network & IT Ops




Figure 5-9. Example graph for the data center deployment scenario

Relationship Queries Strain Traditional Databases



A single query can touch a *lot of data*



Queries can take non-sequential, *arbitrary paths* through data



Real-time queries need speed and *consistent response times*



Queries must *run reliably* with *consistent results*

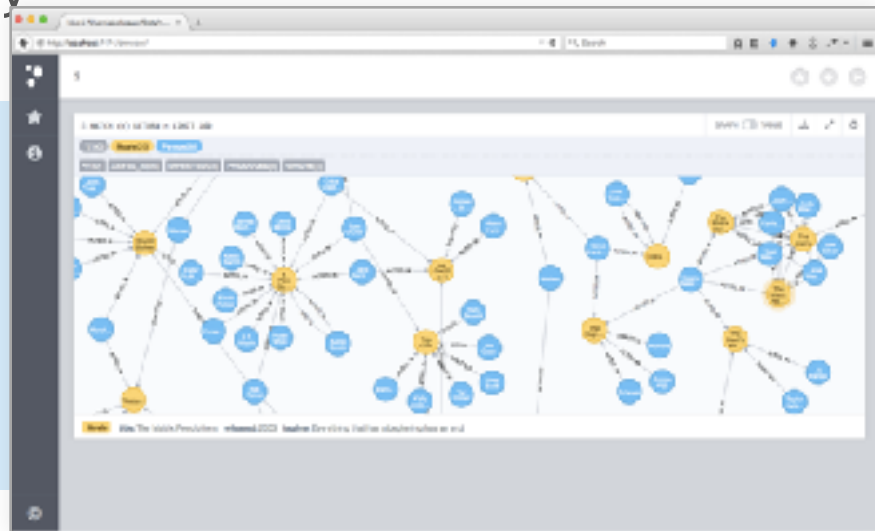
Neo4j - The #1 Platform for Connected Data

Neo4j is an *enterprise-grade native graph database* that enables you to:

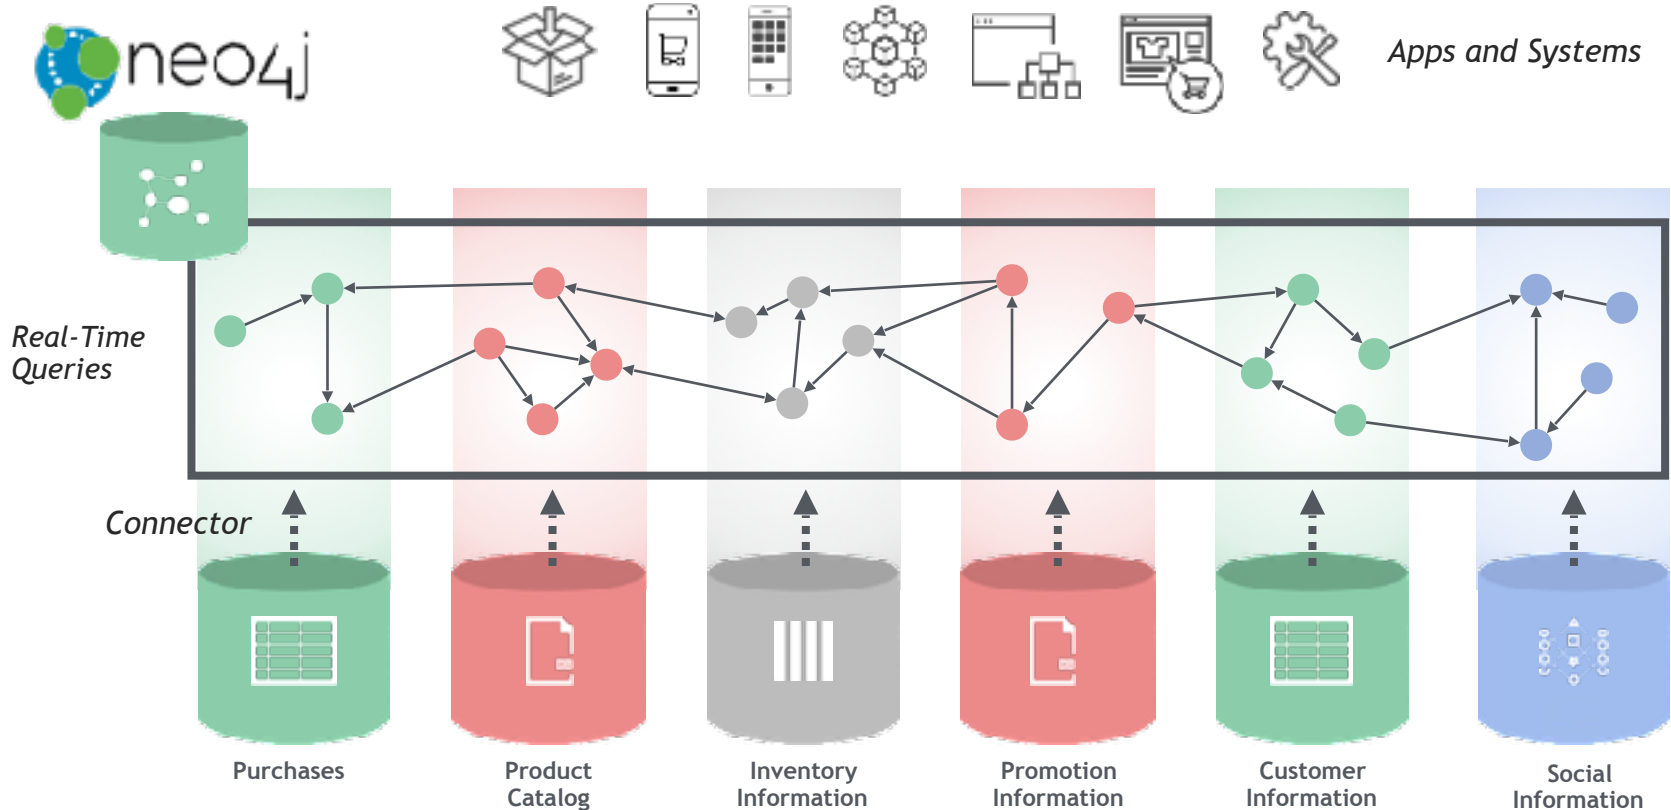
- **Store and query** data relationships
- **Traverse** any levels of depth on real-time
- **Add and connect** new data on the fly

Designed, built and tested *natively* for graphs from the start to ensure:

- Performance
- ACID Transactions
- Agility
- Developer Productivity
- Hardware Efficiency



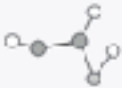









Neo4j – Brings Together and Mobilizes Your Data



Lessons Learned: Ten Year Head Start

Native Connectedness Differentiates Neo4j

	 Native Graph DB	Non-Native Graph DB	RDBMS
Conceive			
Code	Cypher <code>(graphs)-[:are]->(energyhere)</code>	Cypher/Gremlin /Proprietary	SQL
Compute			
Store			

Optimized for graph workloads

NoSQL Databases Are NOT All the Same



MongoDB was clearly the wrong NoSQL database for providing recommendations and was a significant factor in a \$250 million drop in market cap.

‘...we’ve heard one former employee ... describe it as a “technical s**t show.” ’

<https://techcrunch.com/2017/06/09/pandora-raises-480m-from-siriusxm-sells-ticketfly-to-eventbrite-for-200m>

Transformative Customer Experiences

Real-time promotion recommendations

FORBES
50
RETAIL

- Record “Cyber Monday” sales
- About 35M daily transactions
- Each transaction is 3-22 hops
- Queries executed in 4ms or less
- Replaced IBM Websphere commerce



Marriott's Real-time Pricing Engine




- 300M pricing operations per day
- 10x transaction throughput on half the hardware compared to Oracle
- Replaced Oracle database



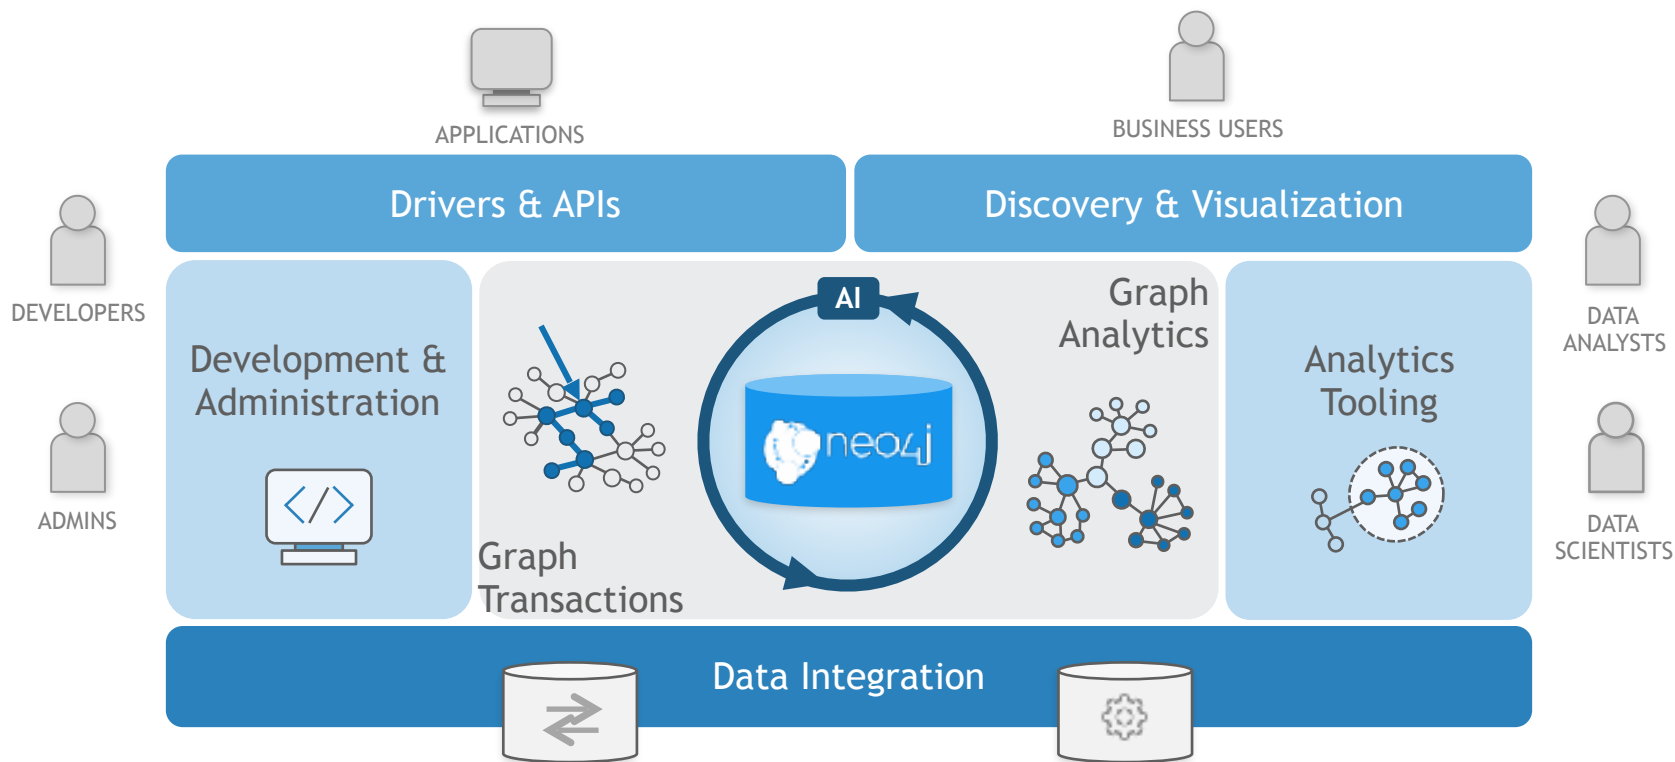
Handling Package Routing in Real-Time



- Large postal service with over 500K employees
- Neo4j routes 7M+ packages daily at peak, with peaks of 5,000+ routing operations per second.



Roadmap - Neo4j Graph Platform

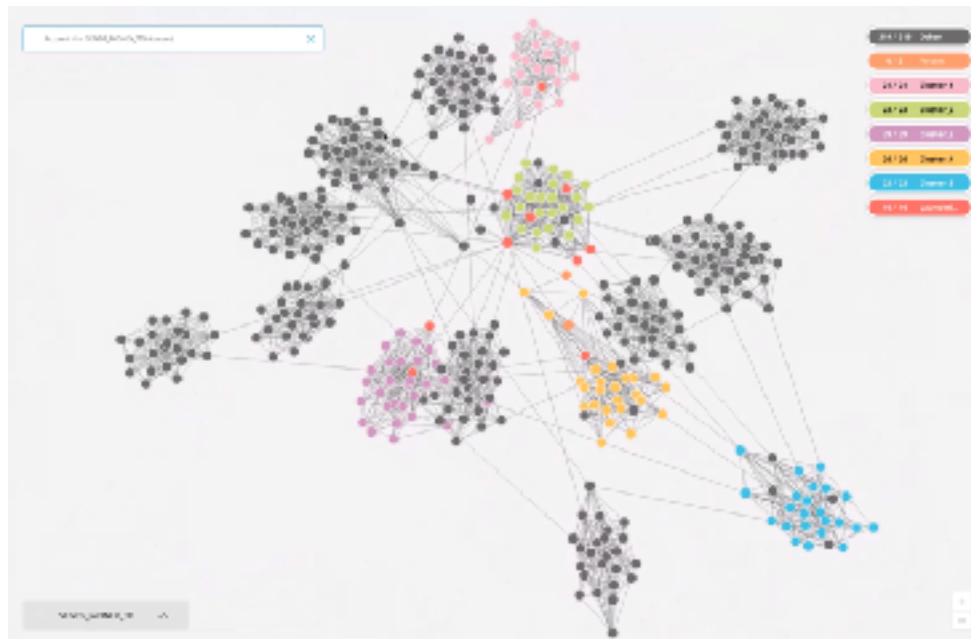


Neo4j Roadmap - Introducing Neo4j Bloom

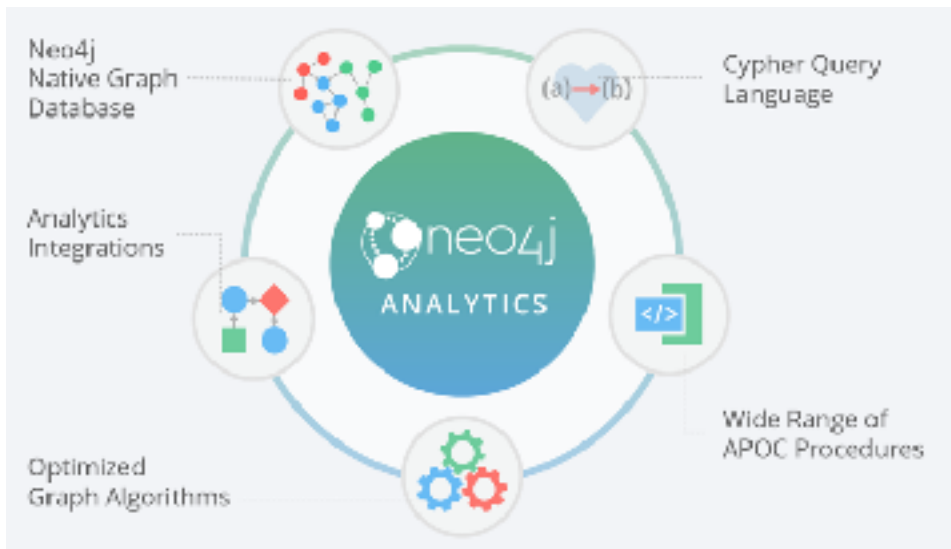
- Graph visualization accessible to the entire enterprise
- Minimal training required to derive insights from connected data using Bloom

Neo4j Bloom Planned 1.0 Features

Graph Perspectives	a business user view of the graph
Graph Visualization	high performance graph layout + rendering
Graph Exploration	navigate through direct graph interaction
Graph Inspection	browseable details of graph entities
Graph Editing	create, duplicate, edit, delete
Graph Search	extensible, idiomatic search phrases



Neo4j Graph Analytics and Algorithms



Graph Algorithms - Insights

- Metrics
- Relevance
- Clustering and Classification
- Dependencies

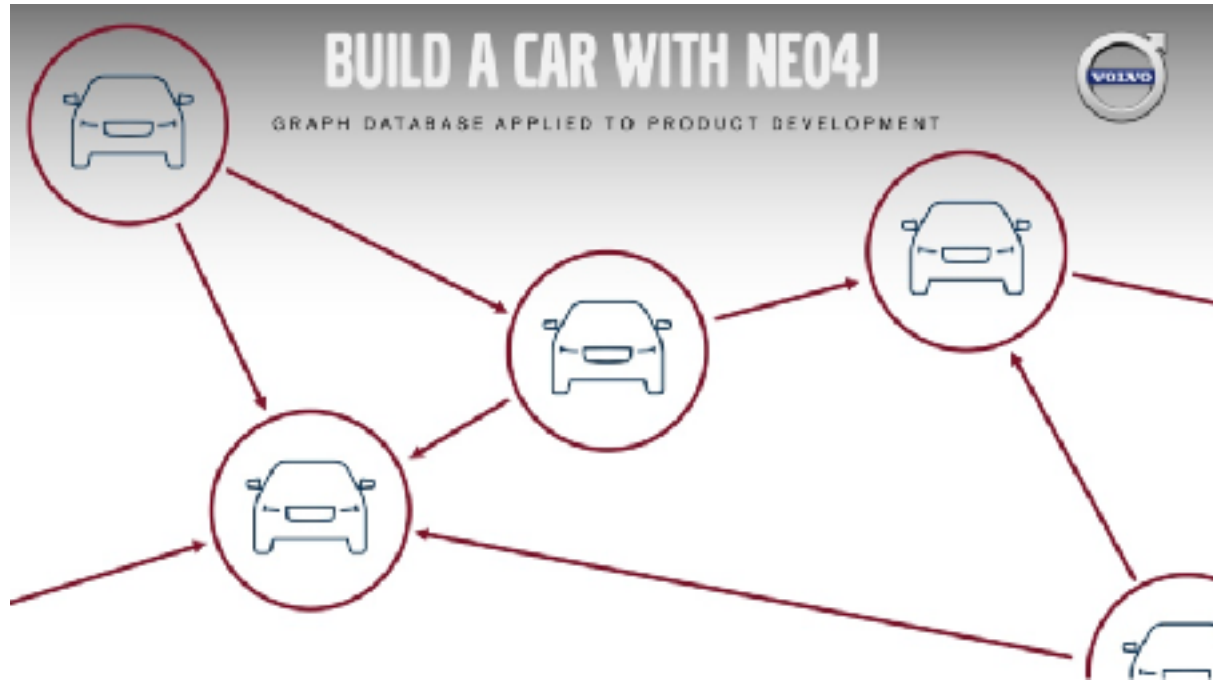




Graphs in Manufacturing



Neo4j at Volvo



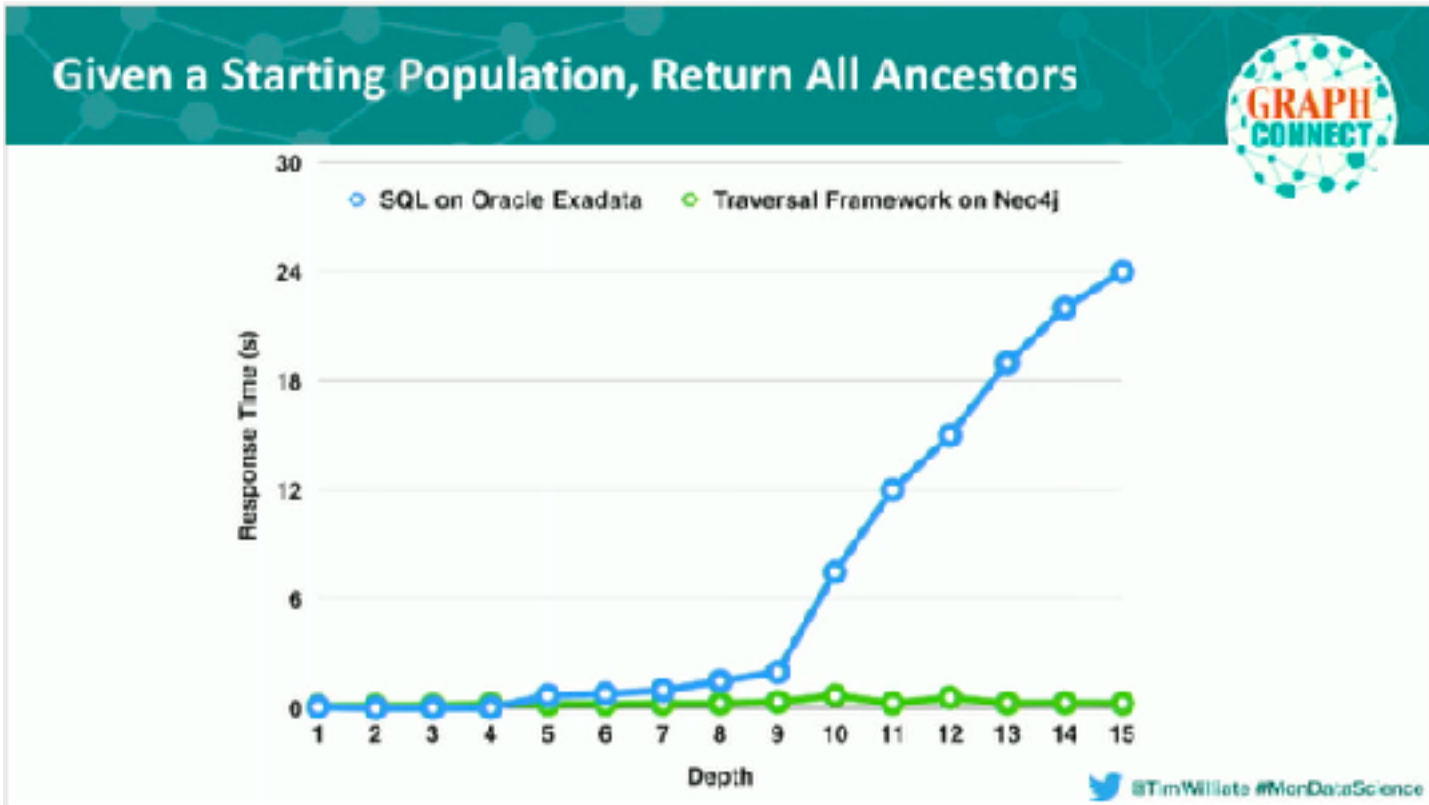
<https://www.slideshare.net/neo4j/volvo-cars-build-a-car-with-neo4j>

Neo4j Supply Chain Gist



<https://neo4j.com/graphgist/supply-chain-management>

Neo4j Replaces Oracle at Monsanto



US Army Supply Logistics




- Forecast the need for replacement parts
- Calculate mean time to failure rates
- Perform multi-dimensional cost comparison and trend analysis
- Inform the Army's budget requirements process
- Answer vital "what-if" questions such as the cost of deploying certain forces and the supporting equipment to a new war zone

<https://neo4j.com/case-studies/us-army/>

Bill of Materials - A Graph Specialty

MRP, ERP, PLM ... Where is DA BOM?



BOM connects all engineering and non-engineering processes

Copyright © Beyond PLM 2015



<https://maxdemarzi.com/2017/11/17/bill-of-materials-in-neo4j/>






Graph Algorithms



Use Graph Algorithms

Extract Structures and Infer Meaning



Use Graph Algorithms



Pathfinding & Search

Finds the optimal path or evaluates route availability and quality



Centrality

Determines the importance of distinct nodes in the network



Community Detection

Evaluates how a group is clustered or partitioned

Algorithms - Pathfinding & Search




- **Single-Source Shortest Path**

- Calculates "shortest" path between a node and all other nodes

- **All-Pairs Shortest Path**

- Finds all shortest paths between all nodes

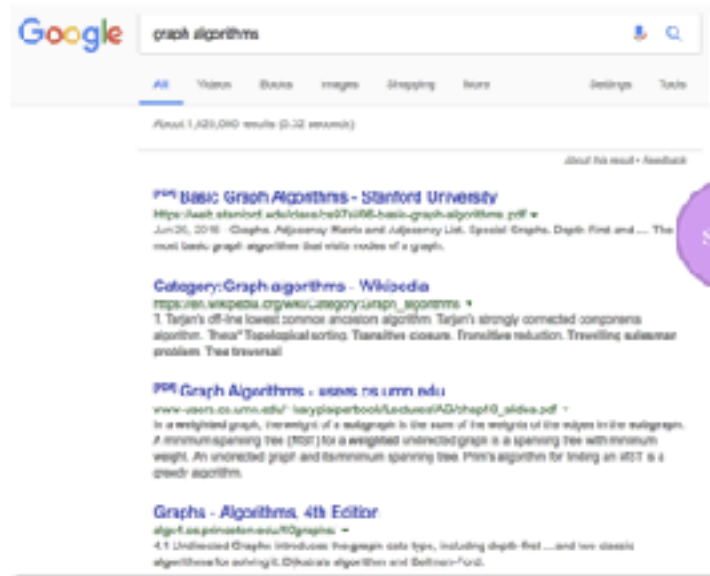


Algorithms - Centralities



- **PageRank**

- Which nodes have the most overall influence



Centrality

Measure of importance

PageRank

- Recursive
- Importance and number of connected nodes

Betweenness Centrality

- Number of shortest paths connecting all pairs in the network

Closeness Centrality


- Inverse of distance to all other nodes in the network



neo4j



Visualization





Technical




The modeling workflow



As an air travel enthusiast
I want to know how airports are connected
So that I can find the busiest ones

Is Airport A connected to Airport B?

Origin	Dest	FlightNum
IAD	TPA	335
IAD	TPA	3231
IND	BWI	448
IND	BWI	3920



(airport) -[:CONNECTED_TO]->(airport)

```
MATCH (origin:Airport {code: "LAX"})  
-[:flight:CONNECTED_TO]->  
  (destination:Airport {code: "LAS"})  
RETURN origin, destination, flight  
LIMIT 10
```

```
CREATE (lax:Airport {code: "LAX"})  
CREATE (las:Airport {code: "LAS"})  
CREATE (las)-[:connection:CONNECTED_TO {  
  airline: "WN",  
  flightNumber: "82",  
  date: "2008-1-3",  
  departure: 1715,  
  arrival: 1820}]->(lax)
```



Cluster Architecture



Raft-based architecture

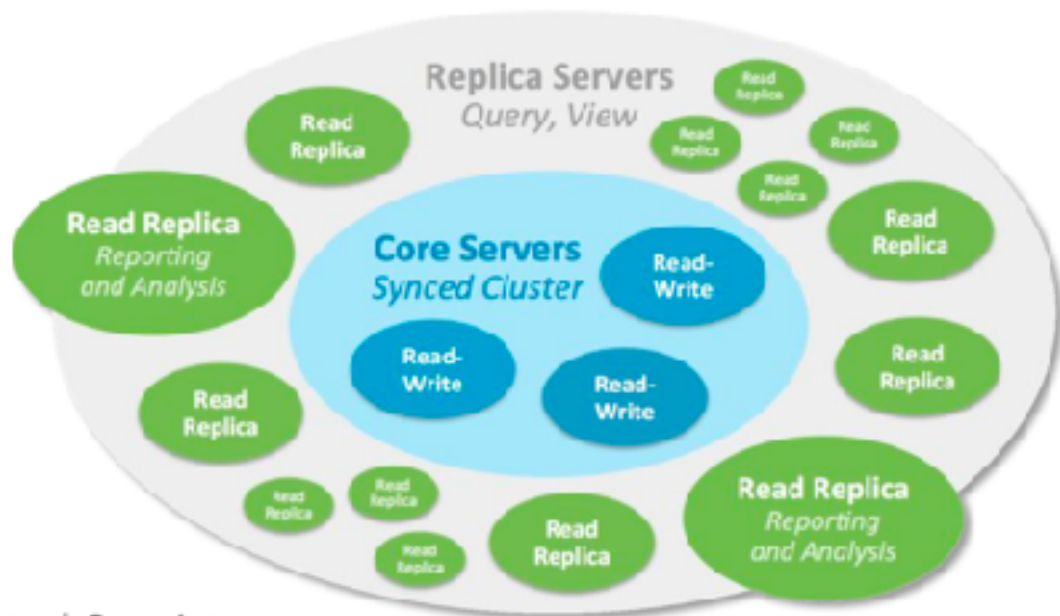
Consensus commits via "Core" servers

Cluster-aware drivers

No need for external load balancer
Stateful, cluster-aware sessions

Causal Consistency

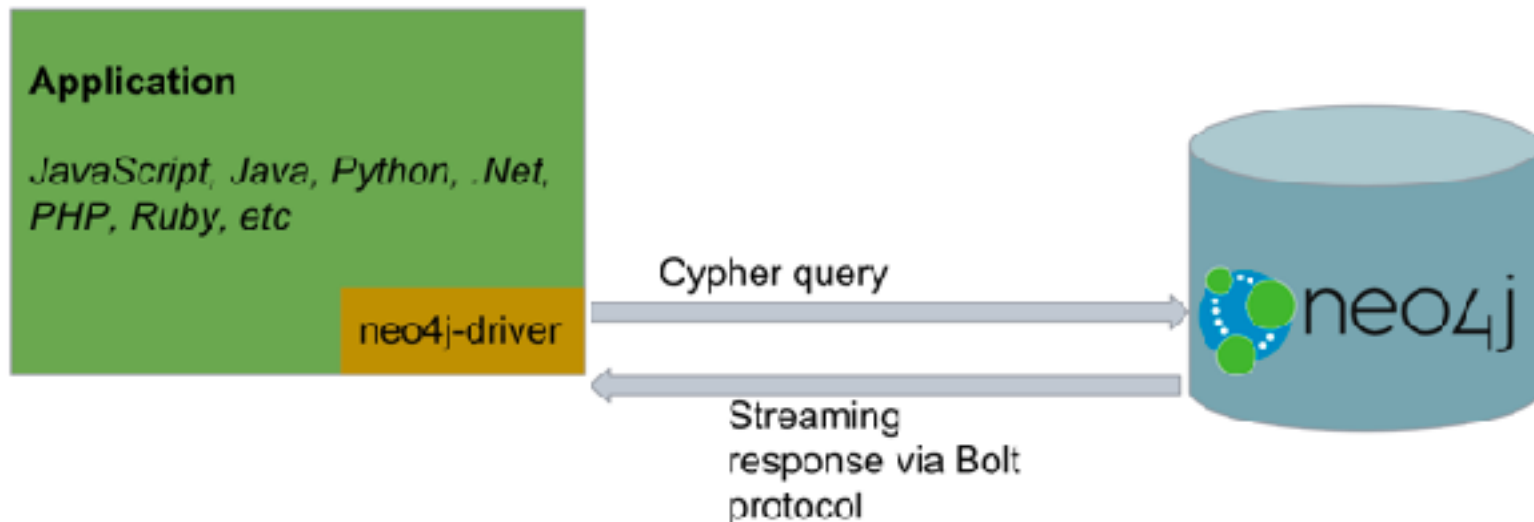
Stronger consistency model than Eventual Consistency



Application Architecture

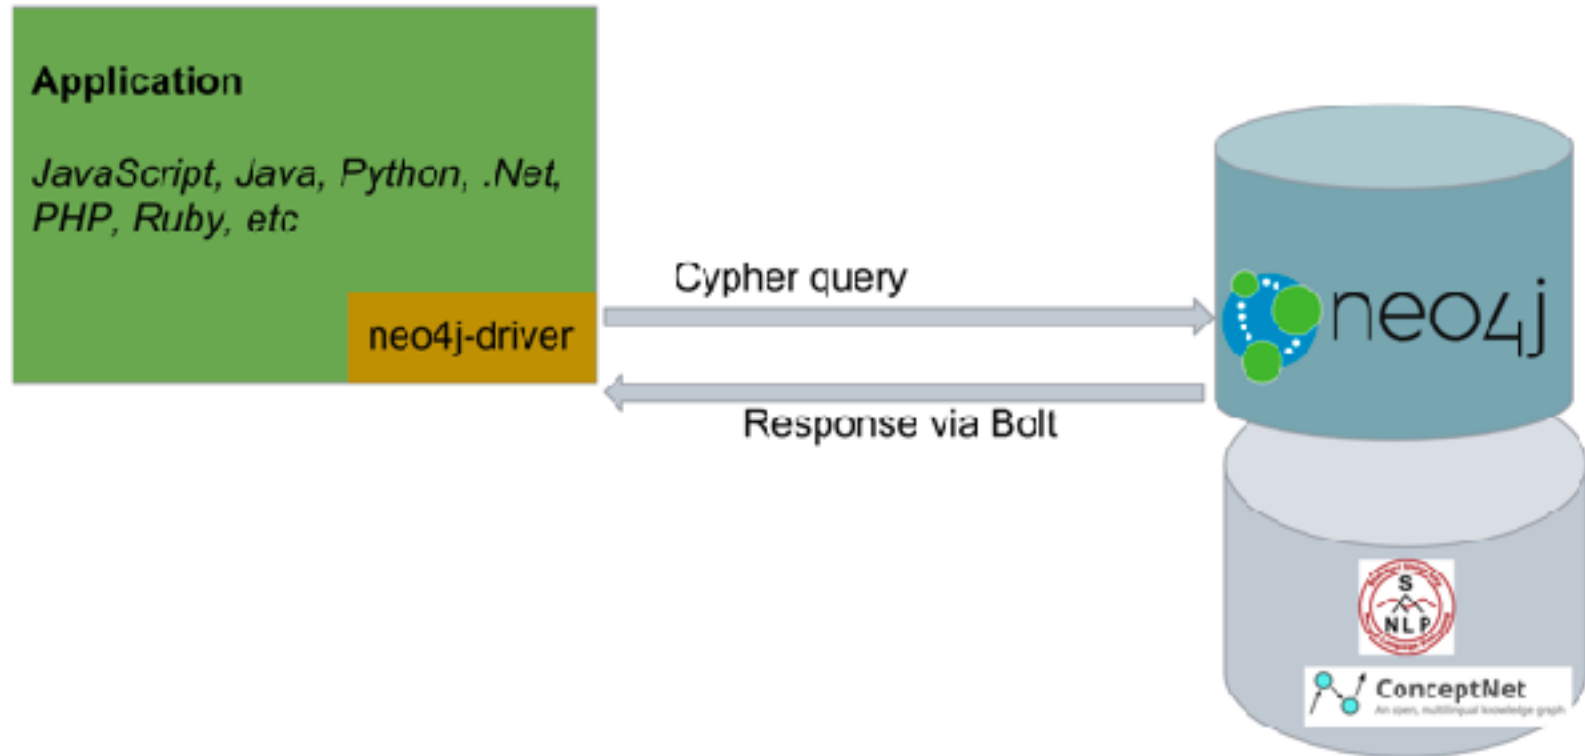


Client driver libraries use Cypher over Bolt protocol



<https://neo4j.com/developer/language-guides/>

Using NLP and User-Defined Functions



Using NLP and User-Defined Functions



GraphAware Natural Language Processing

build **passing** This [Neo4j](#) plugin offers Graph Based Natural Language Processing capabilities.



```
1 MATCH (tw:Tweet {lang: "en"})
2 CALL ga.nlp.annotate({text: tw.text, id: id(tw)})
3 YIELD result
4 MERGE (tw)-[:HAS_ANNOTATED_TEXT]->(result)
5 RETURN count(result)
```

<https://github.com/graphaware/neo4j-nlp>

Neo4j With Python Notebooks



johnnymontana / graph-of-thrones

Watch 4 Star 22 Fork 14

Code Issues 0 Pull requests 1 Projects 0 Wiki Insights

Branch: master graph-of-thrones / network-of-thrones.ipynb Find file Copy path

johnnymontana add explainer images for pagerank and community detection abdcf27 on Aug 8, 2016

1 contributor

764 lines (763 sloc) | 19.9 KB

<> Raw Blame History

Analyzing the Graph of Thrones with Neo4j

```
In [5]: #! pip install py2neo
from py2neo import Graph
graph = Graph()
```

```
In [3]: !pip install py2neo --upgrade

Collecting py2neo
  Downloading py2neo-3.1.0-py3-none-any.whl (145kB)
    100% |#####| 145kB 3.0MB/s
Installing collected packages: py2neo
  Found existing installation: py2neo 2.0.8
  Uninstalling py2neo-2.0.8:
    Successfully uninstalled py2neo-2.0.8
  Successfully installed py2neo-3.1.0
```

Neo4j vs JanusGraph / DSE Graph / Titan



	Neo4J	Titan & Derivatives
ACID Compliance	Yes	Eventually Consistent
Graph Storage	Native Graph	Cassandra
Suitable for Multi-Hops	Yes	No
Language	Open Cypher	Gremlin

OpenCypher vs Apache Gremlin - Simplicity



	Open Cypher	Apache Gremlin
Language Style	Declarative Pattern Matching	Imperative API Style
Ease of Use	Simple to learn	Extremely Complex
Costs to Develop and Maintain	Low	High
Suitable for Power Users	Yes	Are you kidding?

OpenCypher vs Apache Gremlin - Simplicity



CYPHER

```
CALL algo.betweenness.stream('User','MANAGE',{direction:'out'})
```

vs Gremlin

```
gremlin> g.V().as("v"). ❶
  repeat(both().simplePath().as("v").emit(). ❷
    filter(project("x","y","z").by(select(first,"v")). ❸
      by(select(last,"v")).
      by(select(all,"v").count(local)).as("triple").
    coalesce(select("x","y").as("a"). ❹
      select("triples").unfold().as("t").
      select("x","y").where(eq("a")).
      select("t"),
      store("triples")). ❺
    select("z").as("length").
    select("triple").select("z").where(eq("length")). ❻
  select(all,"v").unfold(). ❼
  groupCount().next() ❽
```

<http://tinkerpop.apache.org/docs/current/recipes/>

Cypher for Apache Spark




```
val CITYFRIENDS_NA = SN_NA.cypher(  
  """  
  | MATCH (a:Person)-[:IS_LOCATED_IN]->(city:City)<-[:IS_LOCATED_IN]-(b:Person),  
  |   {a}-[:KNOWS*1..2]->(b)  
  | WHERE city.name = "New_York" OR city.name = "San Francisco"  
  | RETURN GRAPH result OF (a)-[r:SIMILAR_CIRCLE]->(b)  
  """  
  .stripMargin).graphs("result").cache
```

CITYFRIENDS_NA.asZeppelinGraph



Nodes 37: Person

Relationships 99: SIMILAR_CIRCLE





Thank You!

